

Lecture 7.1

No External Regret Algorithm

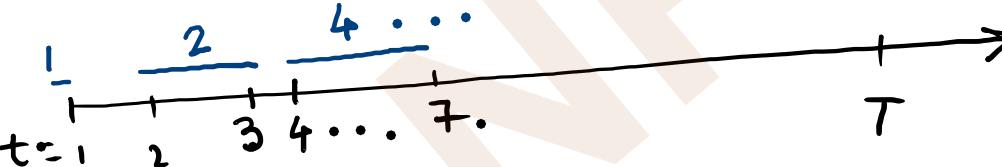
Theorem: $|A| = n$. There exists an algorithm with external regret

(time-averaged) $O\left(\sqrt{\frac{\ln n}{T}}\right)$.

Multiplicative weight algorithm. Set $\varepsilon = \sqrt{\frac{\ln n}{T}}$. The problem with setting $\varepsilon = \sqrt{\frac{\ln n}{T}}$ is that we need to know the time horizon T in advance.

Remark: $|A| = n$. Then there exists an algorithm with/ external time-averaged regret $O\left(\sqrt{\frac{n}{T}}\right)$. Moreover the algorithm does not need to know T a-priori.

Proof:



Group the iterations into epochs. ξ_i , $i = 0, \dots, l$

$$\xi_i = \left[2^i, \min\{2^{i+1} - 1, T\} \right]$$

In the beginning of each epoch, we reset the weight vector of the MW algorithm to the all 1's vector.

In the i -th epoch, we use $\varepsilon = \varepsilon_i = \sqrt{\frac{\ln n}{2^i}}$. Let OPT_i be the pay off of any fixed action in the i -th epoch.

$$\begin{aligned} \text{OPT} - \sum_{i=1}^T r_i &\leq \sum_{i=0}^T \left[\text{OPT}_i - \sum_{t \in \mathcal{E}_i} r_i \right] \\ &\leq \sum_{i=0}^T \left[\varepsilon_i \cdot 2^i + \frac{\ln n}{\varepsilon_i} \right] \end{aligned}$$

$$\begin{aligned}
 &= \sqrt{\ln n} \sum_{i=0}^{\lfloor \frac{T}{2} \rfloor} 2^{\frac{i}{2}} + 1 \\
 &\leq \frac{1}{2} \cdot \sqrt{\ln n} \\
 &\leq 4 \sqrt{T \ln n}
 \end{aligned}$$

□

$$\frac{1}{T} \left(\text{OPT} - \sum_{i=1}^T x_i \right) \leq 4 \sqrt{\frac{\ln n}{T}}$$

Combining Expert Advice: The problem of designing a no-regret algorithm is also called "combining expert

"advice" - the mixed strategy of every iteration is like "combining experts' advices" and the goal is to perform as good as an expert asymptotically.

Combining exploration and exploitation.

Connection between no-external-regret dynamic and

equilibrium concepts:

Let $T = \langle N, (S_i)_{i \in N}, (u_i)_{i \in N} \rangle$ be any normal form game.

Each player runs a no-external regret algorithm.

$$A_i = S_i$$

$$p_t \in \Delta(S_i)$$

$$\pi_i(s_i) = u_i(s_i, p_{-i}) = \sum_{s_{-i} \in S_{-i}} \left[u_i(s_i, s_{-i}) \prod_{j \neq i} p_j(s_j) \right]$$

Theorem: Let $\varepsilon > 0$, players run their no-external-regret algorithm for T iterations such that their time-averaged external regret is at most ε . Define $\sigma_i = \overline{\mathbb{P}}_{i \in N}^T p_i^t$,

$$\sigma = \frac{1}{T} \sum_{i=1}^T \sigma_i. \quad \text{Then } \sigma \text{ is an } \varepsilon\text{-CCE of } T.$$

That is $\mathbb{E}_{s \sim \sigma} [u_i(s)] \geq \mathbb{E}_{s \sim \sigma} [u_i(s'_i, s_{-i})] - \varepsilon \quad \forall i \in N$

$$\forall s'_i \in S_i$$

Lecture 7.2

Theorem: Let $\varepsilon > 0$. Let T be such that the time-averaged external regret of every player is at most ε .

Define $\sigma_t = \prod_{i \in N} p_i^t$ and $\sigma = \frac{1}{T} \sum_{t=1}^T \sigma_t$. Then σ is an ε -CCE.

Proof:

$$\sigma = \frac{1}{T} \sum_{t=1}^T \sigma_t$$

$$\mathbb{E}_{\pi \sim \sigma} [u_i(\pi)] = \frac{1}{T} \sum_{t=1}^T \mathbb{E}_{\pi \sim \sigma_t} [u_i(\pi)]$$

For every $i \in N$, every $\lambda'_i \in S_i$

$$\frac{1}{T} \left[\sum_{t=1}^T \underbrace{\mathbb{E}_{(\lambda_i^t, \lambda_{-i}^t) \sim \sigma_t} [u_i(\lambda'_i, \lambda_{-i}^t)]} - \sum_{t=1}^T \mathbb{E}_{\lambda \sim \sigma} [u_i(\lambda)] \right] \leq \varepsilon$$

$$\underbrace{\mathbb{E}_{(\lambda_i^t, \lambda_{-i}^t) \sim \sigma} [u_i(\lambda'_i, \lambda_{-i}^t)]} - \mathbb{E}_{\lambda \sim \sigma} [u_i(\lambda)] \leq \varepsilon$$

Hence, σ is an ε -CEE.

Q. Does there exist similar natural learning dynamic which enables player to converge to a CE?

Swap-Regret:

"Modification rule" takes the history of the play as input and outputs a sequence of actions for t iterations.
 $(\pi_1, \dots, \pi_t, p_1, \dots, p_t, a_1, \dots, a_t)$ history.

External regret tries to perform as good as any fixed actions; i.e. the modification rules $\{F_a = (\underbrace{a, \dots, a}_t) \mid a \in A\}$

Swap regret tries to perform as good as any modification rule $\{F_{a,b} : a, b \in A, a \neq b\}$

↑
whenever the player played a , it outputs b and vice versa.

No Swap-Regret: A learning algorithm is said to have no swap regret if the time-averaged swap regret

$$\frac{1}{T} \left(\max_{\delta \in \mathcal{F}^{\text{swap}}} \sum_{t=1}^T \sum_{a \in A} p_t(a) \cdot \pi_t(\delta(a)) - \sum_{t=1}^T \sum_{a \in A} p_t(a) \pi_t(a) \right)$$

goes to 0 as T goes to ∞ .

- The connection of no swap-regret algorithm and CE:

Theorem: Let $\varepsilon > 0$. Each player runs a no-swap-regret algorithm.

$$A_i = S_i, \quad \pi_i(s_i) = u_i(s_i, p_{-i}) = \sum_{(s_j)_{j \neq i} \in S_{-i}} u_i(s_i, s_{-i}). p_j(s_j)$$

Let T be such that the time-averaged swap-regret of every player is at most ε . Define $\sigma_t = \frac{1}{T} \sum_{t=1}^T p_t$ and

$$\sigma = \frac{1}{T} \sum_{t=1}^T \sigma_t. \quad \text{Then } \sigma \text{ is an } \varepsilon - CE \text{ of } T.$$

Proof: By definition of σ ,

$$\mathbb{E}_{\lambda \sim \sigma} [u_i(\lambda)] = \frac{1}{T} \sum_{t=1}^T \mathbb{E}_{\lambda \sim \sigma_t} [u_i(\lambda)]$$

For any $i \in N$ and any $s: S_i \rightarrow S_i$

$$\frac{1}{T} \left[\sum_{t=1}^T \sum_{a \in A} p_i^t(a) \pi_t(a) - \sum_{t=1}^T \sum_{a \in A} p_i^t(a) \pi_t(s(a)) \right] \leq \varepsilon$$

$$\Rightarrow \frac{1}{T} \left[\sum_{t=1}^T \mathbb{E}_{\lambda \sim \sigma_t} [u_i(s(\lambda_i), \lambda_{-i})] - \sum_{t=1}^T \mathbb{E}_{\lambda \sim \sigma_t} [u_i(\lambda)] \right] \leq \varepsilon$$

$$\Rightarrow \mathbb{E}_{\lambda \sim \sigma} [u_i(s(\lambda_i), \lambda_{-i})] - \mathbb{E}_{\lambda \sim \sigma} [u_i(\lambda)] \leq \varepsilon$$

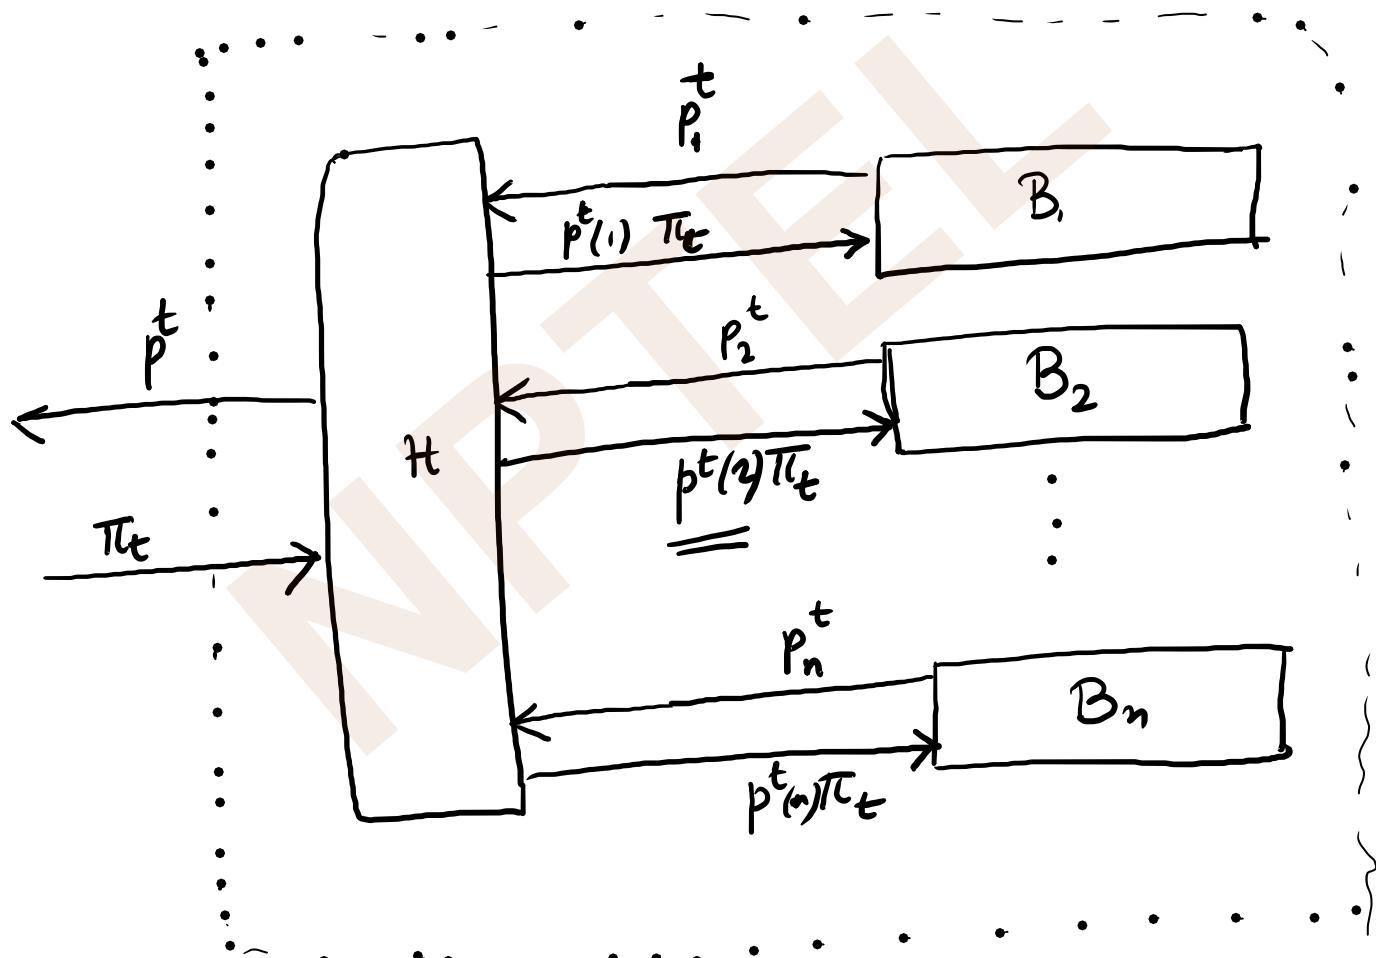
Q. Does there exist a no-swap-regret algorithm?

Black box reduction from no-Swap regret to no external regret:

Theorem: Let $|A|=n$. Suppose there exists an algorithm with time averaged external regret $R(T, n)$. Then there also exists an algorithm with time averaged swap regret $n R(T, n)$. In particular, if there exists a

no external regret algorithm (i.e. $\lim_{T \rightarrow \infty} R(T, n) = 0$), then there also exists a no swap regret algorithm.

Proof: Let the action set of the player (of the no swap regret algorithm) be $A = \{1, \dots, n\}$. Let B be an algorithm with time-averaged external regret $R(T, n)$. Take n copies of B ; let us call them B_1, \dots, B_n . We will build a "master algorithm" H using B_1, \dots, B_t .



The time - averaged expected pay off of the master algorithm

is

$$\frac{1}{T} \sum_{t=1}^T \sum_{i \in A} p^t(i) \cdot \pi_t(i)$$

Let $\delta: A \rightarrow A$ be any switching function. Then the time - averaged expected pay off of the master algorithm (modified by δ) is

$$\frac{1}{T} \sum_{t=1}^T \sum_{i \in A} p^t(i) \pi_t(\delta(i))$$

We need to show,

✓
$$\frac{1}{T} \left[\sum_{t=1}^T \sum_{i \in A} p^t(i) \underline{\pi_t(\delta(i))} - \sum_{t=1}^T \sum_{i \in A} \boxed{p^t(i)} \underline{\pi_t(i)} \right] \leq n \cdot R(T, n)$$

Since the external regret of each B_i is $R(T, n)$, we have the following

✓
$$\frac{1}{T} \left(\sum_{t=1}^T \sum_{i \in A} p^t(i) \underline{\pi_t(\lambda)} - \sum_{t=1}^T \sum_{j \in A} p_i^t(j) p^t(i) \pi_t(j) \right) \leq R(T, n)$$

$\forall i \in A, \lambda \in A$

Put $\gamma = \delta(i)$ in the above inequality.

$$\frac{1}{T} \left(\sum_{t=1}^T p^t(i) \pi_t(\delta(i)) - \sum_{t=1}^T \sum_{j \in A} p_i^t(j) p^t(i) \pi_t(j) \right) \leq R(T, \gamma) \quad \forall i \in A.$$

Adding all the above n inequalities,

$$\frac{1}{T} \left(\sum_{t=1}^T \sum_{i \in A} p^t(i) \cdot \pi_t(\delta(i)) - \sum_{t=1}^T \sum_{i \in A} \sum_{j \in A} p_i^t(j) p^t(i) \pi_t(j) \right) \leq n R(T, \gamma)$$

$$\Rightarrow \frac{1}{T} \left(\underbrace{\sum_{t=1}^T \sum_{i \in A} p^t(i) \pi_t(\delta(i))}_{\text{---}} - \underbrace{\sum_{t=1}^T \sum_{i \in A} \pi_t(i)}_{\text{---}} \right)$$

$$\boxed{\sum_{j \in A} p_j^t(i) p^t(j)} \leq n R(T, \gamma)$$

Set p^t as a solution to the following system of linear equations.

$$p^t(i) = \sum_{j \in A} p_j^t(i) \overline{\pi} \quad \forall i \in A$$

$$\sum_{i \in A} p^t(i) = 1$$

Corollary: $|A|=n$. There exists an algorithm with swap regret $O\left(n \cdot \sqrt{\frac{\ln n}{T}}\right)$.

Lecture 7.4

Price of Anarchy

$$\Gamma = \langle N, (S_i)_{i \in N}, (u_i)_{i \in N} \rangle$$

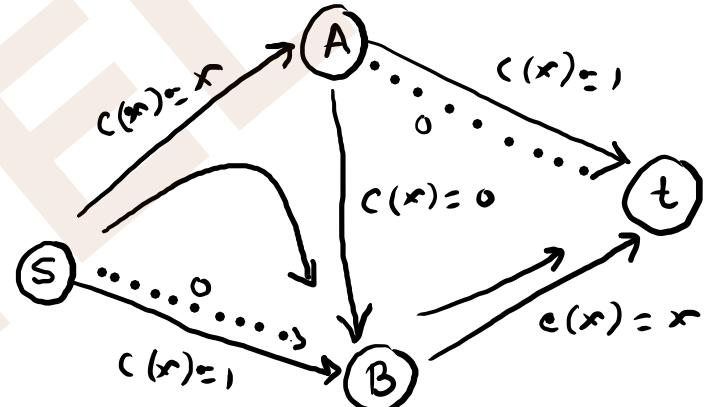
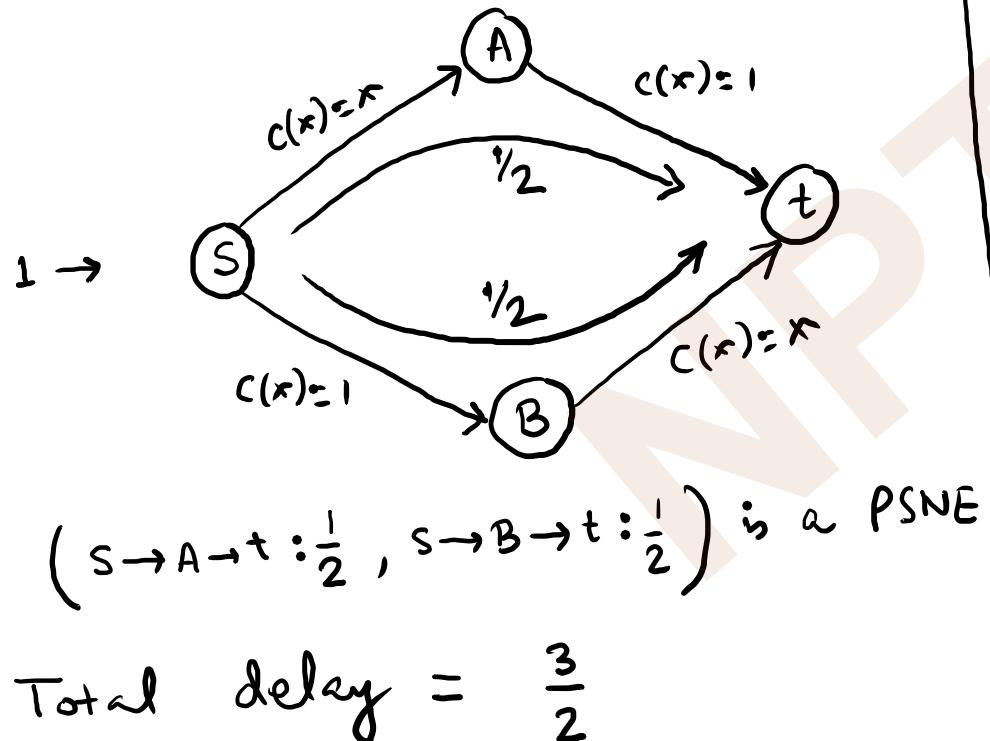
Social welfare function:

$$w: \prod_{i \in N} S_i \rightarrow \mathbb{R}_{>0}$$

Definition (PoA):

$$PoA = \frac{\max_{s^{eq} \in PSNE(\Gamma)} w(s^{eq})}{\max_{s \in \prod_{i \in N} S_i} w(s)}$$

Braess's Paradox



$S \rightarrow A \rightarrow B$ always take a cost of $x (\leq 1)$ which is at most the cost of $S \rightarrow B$

$S \rightarrow A \rightarrow B \rightarrow t$ is a weakly dominant strategy. Thus $(S \rightarrow A \rightarrow B \rightarrow t :)$ is a WDSE.

Total delay = 2

If the social welfare function is $\frac{1}{\text{total delay}}$, then

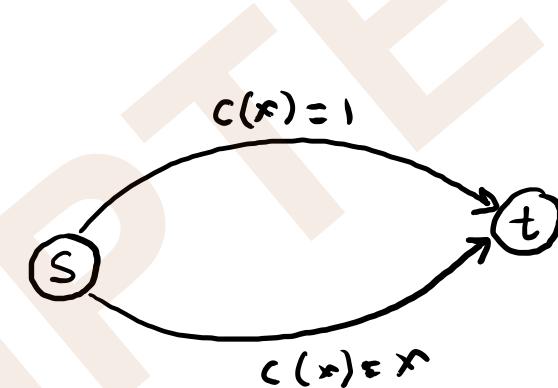
$$\text{PoA} \geq \frac{2}{3/2} = \frac{4}{3}$$

Pigou's Network:

The bottom edge weakly dominates the top edge. So
all traffic using the bottom edge is a WDSE.

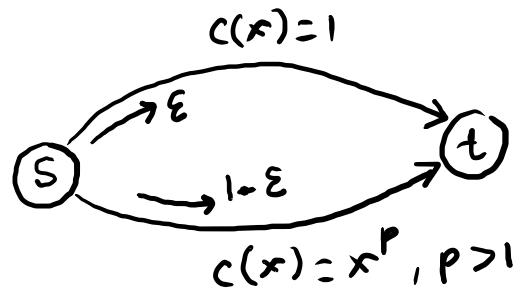
$$\text{Total cost} = 1$$

$\frac{1}{2}$ traffic using the top edge and the other $\frac{1}{2}$ traffic



using the bottom edge incurs a cost of $\frac{1}{2} \cdot 1 + \frac{1}{2} \cdot \frac{1}{2} = \frac{3}{4}$.

$$\text{PoA} \geq \frac{4}{3}$$



All traffic following the bottom path is a WDSE. Total cost is 1.

Total cost is $\varepsilon \cdot 1 + (1 - \varepsilon) \cdot (1 - \varepsilon)^p, \varepsilon > 0$

$$= \varepsilon + (1 - \varepsilon)^{p+1}$$

$$\text{PoA} \geq \frac{1}{\underline{\varepsilon + (1 - \varepsilon)^{p+1}}} \rightarrow \infty \quad \text{as } p \rightarrow \infty$$

As non-linearity in the cost function increases, the PoA also increases.

PoA for Selfish Networks

Pigou's Network:

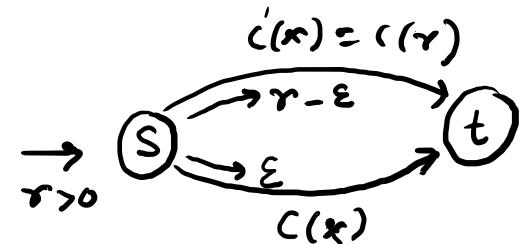
Define:

- It has two vertices: source s and sink t
- edges from s to t .
- Traffic rate is $x (> 0)$
- The cost of one edge is $c(x)$; $c(\cdot)$ is a non-negative, non-decreasing, continuous function

- The cost of the other edge is $c(r)$.

PoA of Pigou's network:

$$\alpha(c) = \sup_{\varepsilon \in [0, r]} \left[\frac{rc(r)}{\varepsilon c(\varepsilon) + (r-\varepsilon)c(r)} \right]$$



Observe that, $\boxed{\varepsilon c(\varepsilon) + (r-\varepsilon)c(r)}$ is non-decreasing in $\varepsilon \in [\gamma, \infty)$ since $c(\cdot)$ is a non-decreasing function.

$$\alpha(c) = \sup_{\varepsilon \geq 0} \frac{rc(r)}{\varepsilon c(\varepsilon) + (r-\varepsilon)c(r)}$$

We let the cost function c belong to a class of functions C .

$$\alpha(C) = \sup_{c \in C} \sup_{r > 0} \alpha(c)$$

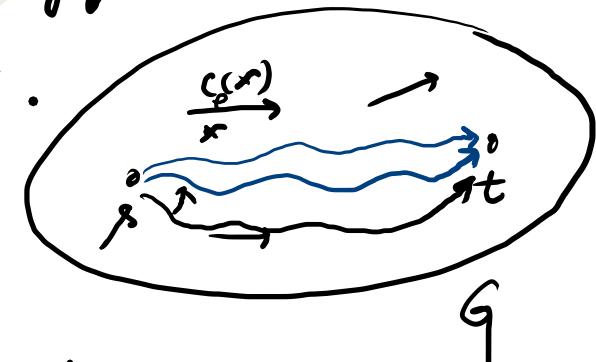
$$= \sup_{c \in C} \sup_{r > 0} \sup_{\varepsilon \geq 0} \left[\frac{rc(r)}{\varepsilon c(\varepsilon) + (r - \varepsilon)c(r)} \right]$$

Theorem: For any class C of cost functions, the PoA of any network with one source and one destination where all the cost functions belong to C , is at most $\alpha(C)$.

proof: Let G be a graph, $f: E[G] \rightarrow \mathbb{R}_{\geq 0}$ be an $s-t$ flow in G . Observe that every strategy profile corresponds to a flow in G .

Observe that a flow f corresponds to a PSNE if and only if only shortest paths carry any flow.

$$f_{\bar{P}} > 0 \Rightarrow \bar{P} \in \operatorname{argmin}_{P \in \mathcal{P}} \left\{ \sum_{e \in P} c_e(f_e) \right\}$$



Let f^* be a flow which minimizes the sum of cost of all the edges.

$$f^* \in \underset{\substack{f \text{ is an } s-t \text{ flow of} \\ \text{value } r}}{\operatorname{argmin}} \left\{ \sum_{e \in E} f_e c_e(f_e) \right\}$$

Since PSNE flow f^{eq} only uses shortest paths according to the cost function $c_e(f_e^{eq})$, $e \in E(g)$, all the $s-t$ paths carrying any flow in f^{eq} must have the same cost.

$$\sum_{P \in \beta} f_P^{eq} \underbrace{c(f_P^{eq})}_{=} = rL \quad , \text{where } L \text{ is the cost of the min-cost path.}$$

—(1)

$$\sum_{P \in \beta} f_P^* \underbrace{c(f_P^{eq})}_{\geq L} \geq L \sum_{P \in \beta} f_P^* = rL \quad —(2)$$

$$(1) - (2)$$

$$\sum_{P \in \beta} (f_P^{eq} - f_P^*) c(f_P^{eq}) \leq 0 \quad —(3)$$

$\alpha(c)$ is supremum over all $c \in C$, $\gamma > 0$, $\varepsilon \geq 0$, we have

the following for any edge $e \in E[G]$,

$$\alpha(c) \geq \frac{f_e^{eq} \cdot c_e(f_e^{eq})}{f_e^* c_e(f_e^*) + (f_e^{eq} - f_e^*) c_e(f_e^{eq})}$$
$$\Rightarrow f_e^* c_e(f_e^*) \geq \frac{f_e^{eq} \cdot c_e(f_e^{eq})}{\alpha(c)} -$$

$$(f_e^{eq} - f_e^*) c_e(f_e^{eq})$$

$$\xrightarrow{\substack{(r =) f_e^{eq} \\ e \quad c_e(f_e^{eq})}}$$

$$\varepsilon = f_e^*$$

$$r = f_e^{eq}$$

$$\Rightarrow \underbrace{\sum_{e \in E} f_e^* c_e(f_e^*)}_{\geq} \geq \frac{1}{\alpha(c)} \underbrace{\sum_{e \in E} f_e^{eq} c_e(f_e^{eq})}_{- \sum_{e \in E} (f_e^{eq} - f_e^*) c_e(f_e^{eq}) \leq 0}$$

$$\Rightarrow c(f^*) \geq \frac{c(f^{eq})}{\alpha(c)}$$

$$\Rightarrow \frac{c(f^{eq})}{c(f^*)} \leq \alpha(c)$$

$$\Rightarrow p_0 A \leq \alpha(c)$$