

Amazon S3 in 3 Minutes: Quick Insights and Practical Examples for Easy Understanding



Nollie Chen · [Follow](#)

4 min read · Nov 7, 2023

Listen

Share


More



How does Amazon S3 work?

Amazon's Simple Storage Service, or [Amazon S3](#), is a key player in the Amazon Web Services (AWS) suite. It's a scalable, highly available, and secure service that excels at providing object storage in the Cloud. This ability enables users to store and retrieve vast amounts of data from anywhere on the web.

The beauty of S3 lies in its flexibility and wide-ranging applications. It can act as a host for static websites, scaling as needed to accommodate traffic. It can serve as a secure storage location for database backups, ensuring data remains safe and easily retrievable. For those involved in data processing, S3 plays a crucial role in storing and retrieving data used in data processing pipelines, making it an invaluable tool in data analytics and machine learning applications.



pc: [cloudfonaut](#)

At its core, Amazon's S3 service revolves around two key concepts: 'Buckets' and 'Objects'. A 'Bucket' is essentially a container or a folder used for storing data, referred to as 'Objects'. Each 'Bucket' must have a globally unique name across all AWS accounts, serving as a unique identifier. Users can create a hierarchical structure within these 'Buckets' using prefixes, simulating a folder-like structure.

'Objects' in these 'Buckets' are the actual data pieces users store in S3. An 'Object' can be any file type — from video files and JSON files to executable JAR files and ZIP files. With the maximum size of an 'Object' being 5TB, S3 offers ample space for large datasets.

Retrieving data from S3 can be done in several ways. While data can be accessed via a URL if its access permissions are set to public, a more common and secure method is to **retrieve data programmatically using AWS SDKs**. For instance, boto3, the AWS SDK for Python, provides an interface to Amazon S3.

Also, to help users manage costs and optimize for specific use cases, Amazon S3 offers various storage classes. These include the Standard Tier for frequently accessed data, the Intelligent Tier for data with changing access patterns, the Infrequent Access tier for less frequently accessed data, and the Glacier tiers for long-term archival. Each tier has different pricing and performance characteristics, allowing users to choose the one that best fits their use case.

To demonstrate how one can interact with S3, consider the following JavaScript code snippets using the AWS SDK. This sample code illustrates how to perform common operations such as uploading a file, retrieving a file, and deleting a file from an S3 bucket. It provides a practical example of how to use the 'Upload', 'getObject', and 'deleteObject' methods respectively.

```

// Sample code for S3 Operations
const {
  Upload
} = require("@aws-sdk/lib-storage");

const {
  S3
} = require("@aws-sdk/client-s3");

// dotenv helps manage environment variables
require('dotenv').config();

const fs = require('fs');

// The name of the bucket that you have created
const BUCKET_NAME = 'nameitasyouwant';

// we load credentials from the .env file
const s3 = new S3({
  credentials: {
    accessKeyId: process.env.ID,
    secretAccessKey: process.env.SECRET
  },
  region: 'us-east-1',
});

// upload a file
const uploadFile = async (fileContent, fileName) => {
  console.log('fileName', fileName);
  // Setting up S3 upload parameters
  const params = {
    Bucket: BUCKET_NAME,
    Key: fileName, // File name we want to upload
    Body: fileContent // the buffer
  };

  // Uploading files to the bucket

  const data = await new Upload({
    client: s3,
    params
  }).done();
  console.log(`File uploaded successfully. ${data.Location}`);
  // return the URL of the object on S3
  return data.Location;
};

```

```

// retrieve a file
const retrieveFile = (fileName) => {
  // Setting up S3 read parameters
  const params = {
    Bucket: BUCKET_NAME,
    Key: fileName, // File name we want to retrieve
  };

  // download file from the bucket
  s3.getObject(params, function(err, data) {
    if (err) {
      throw err;
    }
    console.log(`File downloaded successfully. ${data.Body}`);
    // do something with the file
    const fStream = fs.createWriteStream(`${fileName}`);
    fStream.write(data.Body);
    fStream.end();
    // return data
    return data.Body;
  });
};

// delete a file
const deleteFile = (fileName) => {
  // Setting up S3 delete parameters
  const params = {
    Bucket: BUCKET_NAME,
    Key: fileName, // File name we want to delete
  };

  // download file from the bucket
  s3.deleteObject(params, function(err, data) {
    if (err) {
      // throw err;
      return false;
    }
    console.log(`File deleted successfully. ${data}`);
    return true;
  });
};

module.exports = {uploadFile, retrieveFile, deleteFile}

```

In summary, Amazon S3 is an incredibly powerful and flexible storage service within the AWS ecosystem. Its durability, scalability, and high availability make it an

excellent choice for a wide range of applications and use cases. Whether you're hosting a website, backing up data, or executing complex data processing tasks, Amazon S3 has the features and capabilities to meet your storage needs.

Relevant Resources: [AWS Amazon S3 Getting Started](#)

Aws S3

AWS

JavaScript



Follow

Written by Nollie Chen

601 Followers · 22 Following

SDE Intern @AWS | @UPenn | CS gradute | nolliechy@gmail.com | ig:alconollie | linkedin: HuiYu(Nollie) Chen

No responses yet



What are your thoughts?

Respond

More from Nollie Chen